



# **Which whale is it, anyway?**


Face recognition for right whales using deep learning

Robert Bogucki & Maciek Klimek


# The challenge




# Kaggle competition

#	Δrank	Team Name <small>‡ model uploaded * in the money</small>	Score <small>?</small>	Entries	Last Submission UTC (Best – Last Submission)
1	—	deepsense.io <small>‡ *</small>	0.59600	110	Thu, 07 Jan 2016 23:58:03 (-0.2h)
2	—	 felixlaumon <small>‡ *</small>	1.07585	35	Thu, 07 Jan 2016 14:11:52 (-1.6h)
3	—	SKE <small>‡ *</small>	1.14982	100	Thu, 07 Jan 2016 23:51:28 (-0.2h)
4	—	threedB	1.33648	37	Thu, 07 Jan 2016 23:46:31
5	—	AbdulWahab	1.46909	7	Thu, 07 Jan 2016 23:42:40
6	—	Tsakalis Kostas	1.51900	60	Thu, 07 Jan 2016 22:52:43 (-11h)
7	—	bawdyb .	1.55823	49	Thu, 07 Jan 2016 23:30:34 (-23.8h)
8	—	Left Whales <small>‡</small>	1.75764	32	Thu, 07 Jan 2016 23:54:35 (-0h)
9	—	Anil Thomas	1.80178	19	Thu, 07 Jan 2016 06:10:13 (-38.3h)
10	—	Doug Koch	2.13797	85	Thu, 07 Jan 2016 21:09:17 (-9h)

# The true face of a right whale




# The true face of a right whale





# Why is it difficult?




A non-random sample from the dataset

# Why is it difficult?



Not very cooperative whales

# May I see your passport, please?





# Outline of the solution


1




2




3




4



# Where's Your Head At?

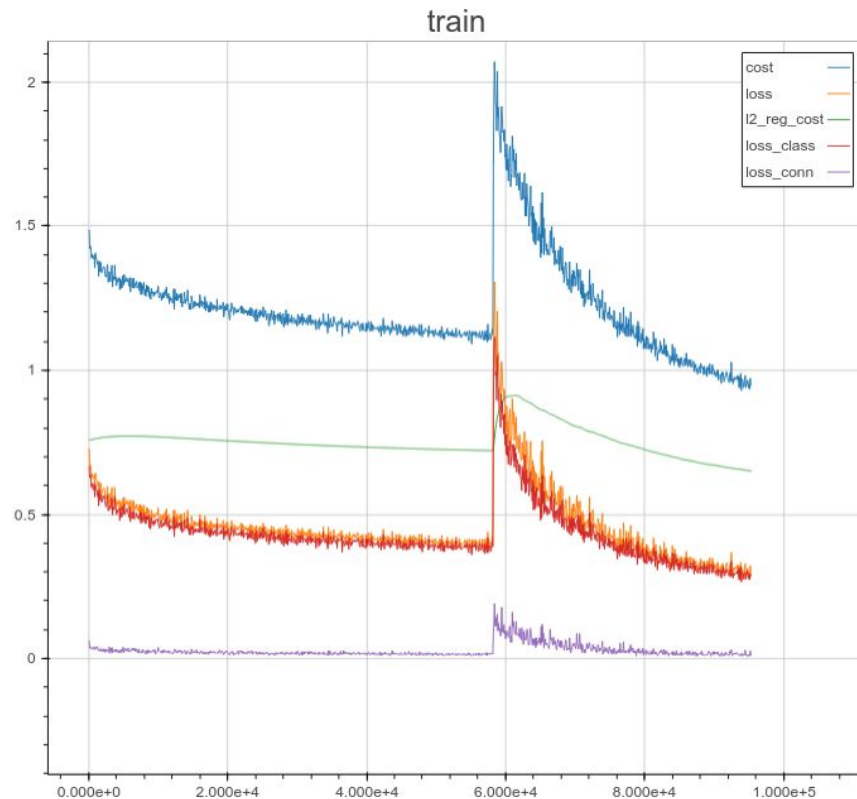



# May I see your passport, please?




# Final classifier

- 87% accuracy, 95% top5 accuracy
- always ensemble
- data augmentation everywhere
- “kicking” the learning rate
- auxiliary targets
- reusing validation set
- calibrating probabilities






# How we got there?



# How we got there?

- Naive classification
- We have to label!
- Sloth




# How to use the annotations?

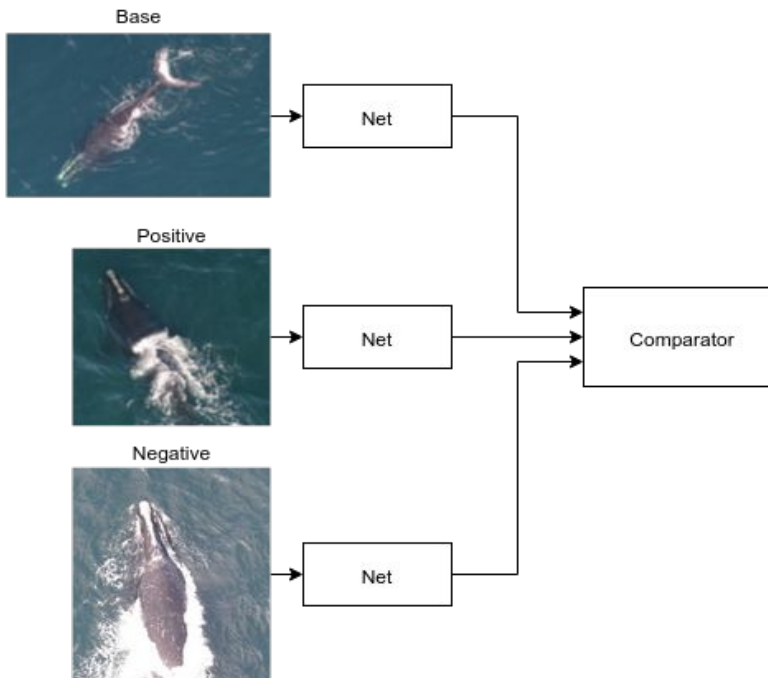
- Positive/negative sliding window
- Dividing the output into bins
- Bonnet-tip and blowhole annotations

# Other things we tried, but...

## Transformer Networks

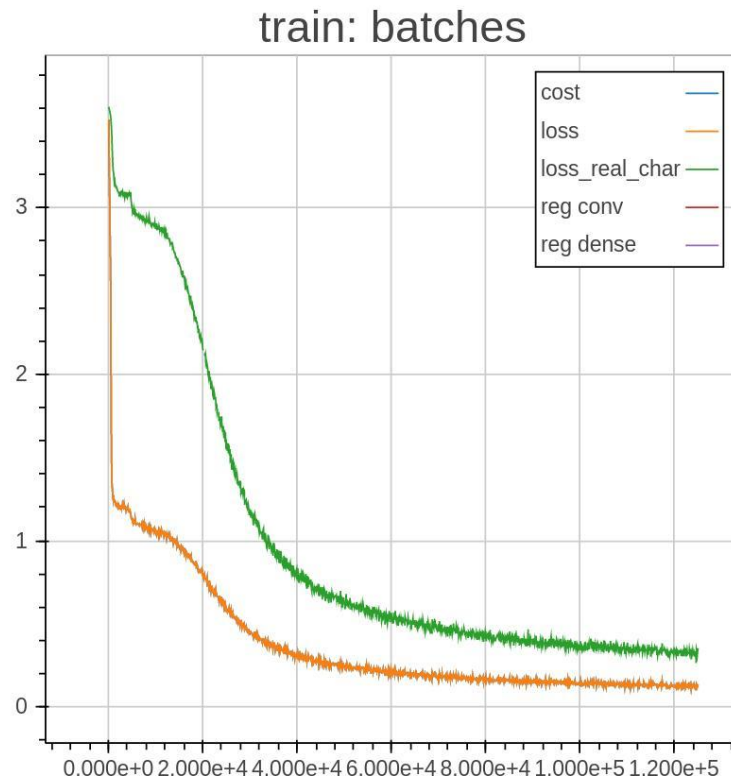


## Triplet training





- Status of the training, plots
- Training many models in parallel
- History and reproducibility
- Sharing results with others



# Keen to learn more?

- [robert@deepsense.io](mailto:robert@deepsense.io)
- [maciej.klimek@deepsense.io](mailto:maciej.klimek@deepsense.io)
- <http://deepsense.io/deep-learning-right-whale-recognition-kaggle/>
- We are hiring!
- Ask questions!

